

Anatomy of an R Help File

Elise Hellwig

Table of contents

1	Overview	3
1.1	Description	3
1.2	Learning Goals	3
1.3	Setup	3
2	Getting Help in R	4
2.1	Looking for Help	4
2.2	Opening R Documentation	4
2.3	Definitions	6
2.3.1	Functions	6
2.3.2	Arguments	6
2.3.3	Objects	6
2.3.4	Classes	7
2.3.5	Methods	7
3	Reading R Help Files	10
3.1	Sections of an R Help File	10
3.2	Title	10
3.3	Description	11
3.4	Usage	11
3.5	Arguments	13
3.5.1	The ... Argument	13
3.6	Details	16
3.7	Value	16
3.8	Note	17
3.9	Author(s)	17
3.10	References	17
3.11	See Also	17
3.12	Examples	18
3.13	Further Reading	19

1 Overview

1.1 Description

The R programming language has a vast and sprawling documentation library, and a lot of it is in a standardized format: the R help file. Unfortunately, R help files can be difficult to read at the best of times, and almost impossible to understand for those new to coding. This workshop will explain each section of the R help file, how to read them, and when they are the most useful. It will also introduce you a number of ways to search R documentation, both from R and on the internet.

Workshop site: https://d-rug.github.io/help_file_anatomy/

1.2 Learning Goals

- Know where to look for R documentation depending on what type of question you need answered
- Understand what information each section of a R help file contains
- Know which help file sections to look at depending on what type of problem you have

1.3 Setup

```
if (!require('sf')) install.packages('sf')
if (!require('forcats')) install.packages('forcats')
```

2 Getting Help in R

2.1 Looking for Help

When programming, it is not a matter of if you get stuck, or even when you get stuck, but how much time will you spend stuck. Some problems can be resolved in minutes where as others may last weeks or even months. Thankfully R has a wide variety of tools to help you get unstuck, from built-in documentation, to mailing lists, to help groups like the Davis R Users Group.

If you have ever asked for coding help, you may have gotten the response “Have you checked the documentation?”. This is the programming equivalent of asking if you have read the manual. And like many manuals, R documentation often seems like it was written for someone who already knows what they are doing. This is all well and good if you do know what you are doing. However, it doesn’t help much if the documentation itself is a source of confusion.

The basic unit of documentation in R is the R help file. While help files may be difficult to decipher at first, they are all difficult in the same way. That means once you understand one of them, understanding the rest is much easier. Most R help files focus on a particular function, like `summary()`, which summarizes various types of data in R. However, many packages and built-in data sets like `mtcars` also have help pages.

2.2 Opening R Documentation

You can access R help files in several ways. If you know the name of the function or package you are looking for information about, you can use the `?` operator in the console with the name of the function to open the R help file. This will only open documentation for functions and packages you currently have **loaded** into R. This is also equivalent to using the `help()` function.

```
#these lines of code are equivalent
?summary

help('summary')
```

If you are using RStudio this will bring up documentation for that function in help window. If you use the GUI, it will open a separate help window. If you are running R in bash or zsh, R will open the help file in your default text editor.

Not all packages have help files the way functions do, but it is getting more common. To load the help file for a package, you need to add `-package` onto the name of the package. Since R does not view “sf-package” as a single word, you need to surround it in tick marks (`'`).

```
#load documentation for the sf package
?'sf-package'
```

If you are unsure of the function name, you can use the `??` operator to search all of your local R documentation for a term or set of terms. For multiple search terms, enclose them in tick marks. The `??` operator is equivalent to the function `help.search()`, but for `help.search` you enclose your search term in quotes. The `help.search` function also gives you more control over what you are searching for.

```
??`linear model`

#look for linear model only in the title of the help file
help.search('linear model', fields='title')
```

This will search documentation for all packages you have **installed**. It will then provide a list links to documentation pages you can use to find the function that works best. Functions are labeled using the schema `[PACKAGE NAME]::[FUNCTION NAME]` so you can tell which package to load to access function you are interested in.

If you want to search documentation for packages that you haven't installed yet, you have a couple of options. If you are only looking for packages that are available through [CRAN](#) or [bioconductor](#), you can use the `RSiteSearch()` function, or [rdocumentation.org](#) if you prefer a web interface. CRAN itself contains a lot of documentation for each package on the package's page (ex. [dplyr](#)), but the information isn't searchable from the website and you have to know what package you want information on. If you also want to be able to search packages hosted on R-forge and Github in addition to CRAN and bioconductor, you can use [rdr.io](#).

Documentation Duplication

Sometimes you will have functions from different packages that share the same name. For example, the `stats` package and the `dplyr` package both have functions called `filter`, but one is used for filtering time series and the other subsets `data.frames` by row. If you use the `?` operator to open the help file for one of these functions, a page will open with links to different documentation pages based on the package you are using. Be sure to click on the link that corresponds to the function you have questions about.

2.3 Definitions

R documentation uses a couple of terms that are sometimes glossed over at introductory levels. However, without them, many R help files are nigh impossible to read.

2.3.1 Functions

A function is a pre-programmed set of instructions that performs a certain task. For example, the `sqrt` function takes the square root of a given number or numbers. The `length` function tells you the number of elements in a vector or list. For more information on functions, see DataLab's [Calling Functions](#) and [Functions](#) sections of the R Basics reader.

2.3.2 Arguments

An argument is an input to a function, the information the function needs in order to work. Not all functions require arguments, but most do. Almost all arguments in R have names. Ideally argument name would tell you something about what the argument does, but that is not always the case. The name of the first three arguments to the `lm` function are `formula`, `data`, and `subset`, which are relatively informative. On the other hand, `sqrt()`'s singular argument `x` doesn't give us very much information, which is why the Arguments section of help files are is important. There are a few functions with unnamed arguments, but we will discuss that later.

2.3.3 Objects

Everything you can interact with in R is an object. Whenever you create a variable, you create an object. All of the objects you create will appear in the Environment tab in RStudio. If you aren't using RStudio, you can also get a list of objects you created using `ls()`.

```
x = 5  
  
ls()
```

```
[1] "x"
```

There are many more objects that R defines internally. R includes many built-in data sets (ex. `mtcars`), all of which are objects. Every function you use is also an object.

2.3.4 Classes

A class is the blueprint for the structure of an object. It defines how you can interact with the object and what types of information it can contain. To determine the class of an object, you can use the function `class()`. Common examples of classes include integers, characters, data.frames, lists, and functions. For more information on classes, see the [Data Types and Classes](#) section of DataLab's R Basics reader.

2.3.5 Methods

A method is a function that is defined for a particular class of object. For example, the function `st_area` calculates the area of a set of shapes. However, it only works for objects of the `sf` class, a type of object that stores geographical information. R assumes that the first argument you provide to `st_area` has the class `sf`, and if it doesn't you will get an error message.

Some functions may have methods defined for many different classes of objects. For example, you can use the summary function to summarize a vector of numbers or the columns of a data.frame. The internal code for the function automatically determines the class of the input data and applies the correct method for you.

```
summary(1:20)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
1.00	5.75	10.50	10.50	15.25	20.00

```
summary(iris)
```

Sepal.Length		Sepal.Width		Petal.Length		Petal.Width	
Min.	:4.300	Min.	:2.000	Min.	:1.000	Min.	:0.100
1st Qu.	:5.100	1st Qu.	:2.800	1st Qu.	:1.600	1st Qu.	:0.300
Median	:5.800	Median	:3.000	Median	:4.350	Median	:1.300
Mean	:5.843	Mean	:3.057	Mean	:3.758	Mean	:1.199
3rd Qu.	:6.400	3rd Qu.	:3.300	3rd Qu.	:5.100	3rd Qu.	:1.800
Max.	:7.900	Max.	:4.400	Max.	:6.900	Max.	:2.500
Species							
setosa	:50						
versicolor	:50						
virginica	:50						

You can use the `methods` function to see what classes of objects have methods for a given function.

```
methods(summary)
```

```
[1] summary,ANY-method          summary,DBIObject-method
[3] summary.aov                 summary.aovlist*
[5] summary.aspell*             summary.check_packages_in_dir*
[7] summary.connection          summary.data.frame
[9] summary.Date                 summary.default
[11] summary.ecdf*               summary.factor
[13] summary.glm                  summary.infl*
[15] summary.lca*                 summary.lm
[17] summary.loess*              summary.manova
[19] summary.matrix              summary.mlm*
[21] summary.nls*                 summary.packageStatus*
[23] summary.POSIXct             summary.POSIXlt
[25] summary.ppr*                 summary.pr_DB*
[27] summary.prcomp*             summary.princomp*
[29] summary.proc_time           summary.proxy_registry*
[31] summary.rlang_error*        summary.rlang_message*
[33] summary.rlang_trace*        summary.rlang_warning*
[35] summary.rlang::list_of_conditions* summary.sfc*
[37] summary.srcfile              summary.scref
[39] summary.stepfun              summary.stl*
[41] summary.svm*                 summary.table
[43] summary.tukeysmooth*         summary.tune*
[45] summary.units*              summary.vctrs_sclr*
[47] summary.vctrs_vctr*          summary.warnings
see '?methods' for accessing help and source code
```

To look up documentation for a specific method, use the schema `[FUNCTION NAME].[CLASS NAME]`. For example, if we want to look up documentation for the `summary` method for the class `lm`, we would use the code `?summary.lm`.

i Advanced Topic

This system of objects, classes, and methods is called Object-Oriented Programming (OOP). Object oriented programming is not unique to R, but R does it a little differently. If you want a deeper understanding of methods in R, the [S3 section](#) of the Intermediate R reader describes the most widely used system of OOP in R. Additionally, [Advanced R](#) by Hadley Wickham provides a more complete description of the various OOP systems

available in R as well as their trade-offs.

3 Reading R Help Files

3.1 Sections of an R Help File

Each R help file has a standard set of sections. Not every help file will have every section, but all help files will have the starred sections in the list below. The sections of an R help file in order are:

- [Title*](#)
- [Description*](#)
- [Usage*](#)
- [Arguments*](#)
- [Details](#)
- [Value](#)
- [Note](#)
- [References](#)
- [Author\(s\)](#)
- [See Also](#)
- [Examples*](#)

3.2 Title

This is the section at the top of the help file and tells you the name of the documentation page you are looking at as well as the package the documentation file is from (enclosed in `{}`). It also displays a very brief summary of what the function does in large lettering. While the title is not something I use regularly, I do occasionally need to figure out what package a particular function lives in and the title is a great place to do that.

If the package that contains the function is the “base” package, that means it is a part of “base” R, or the the set of functions and classes available to you whenever you open an R session. The packages `stats`, `graphics`, `utils`, `datasets`, `methods`, and `grDevices` are also always loaded whenever you open an R session, so you don’t need to call them with `require` or `library`.

3.3 Description

The Description section gives a more detailed description of what the function or package does. It is not generally helpful for troubleshooting. However, it is useful when trying to determine whether or not a function would suit your needs for a particular task.

3.4 Usage

The Usage section is the first regularly useful section, and one of the sections I commonly look at when troubleshooting. It is also probably the most difficult to interpret. In general, it gives you information about how to *use* the function.

Every Usage section will have at least one example function call with the named arguments listed out. Arguments with an = after them have a default value. This means you don't necessarily need to provide a value for that argument when calling the function. However it is important to check that you are okay with whatever the default value is. If there is no = after the argument, it is a required argument and you will need to provide it each time you call the function. The Usage section does not describe what each argument does. For that, you need to go to the Arguments section, which we will cover next. In general, the Usage and Arguments sections are closely links and I often find myself jumping back and forth between them when trying to understand what a function is doing.

Sometimes developers document multiple functions in a single help file. This is the case for `read.table` and `read.csv`, which share a documentation page. Even though two functions share a documentation page, that does not mean they share all the same arguments or default values. For example, the default value for the header argument is `FALSE` for `read.table` but `TRUE` for `read.csv`.

If there is more than one example function call of the same function in the Usage section, that generally means that there is more than one version of the function (method). Which method R uses depends on the object class of the first argument passed to the function. When there is more than one method listed in a given documentation file, the file will tell you which class of object that particular method applies to using a comment just above the function call with language like “method for class X” or “method for signature X”. This tells you that the arguments listed inside that call are the arguments you need to run that method if your first argument has a class of X, as well as the default values for that method.

For example, if we summarize survey data from the [General Social Survey](#) as a `data.frame`, R displays a maximum of 7 different values for categorical variables.

```
library(forcats)

summary(gss_cat)
```

year	marital	age	race
Min. :2000	No answer : 17	Min. :18.00	Other : 1959
1st Qu.:2002	Never married: 5416	1st Qu.:33.00	Black : 3129
Median :2006	Separated : 743	Median :46.00	White :16395
Mean :2007	Divorced : 3383	Mean :47.18	Not applicable: 0
3rd Qu.:2010	Widowed : 1807	3rd Qu.:59.00	
Max. :2014	Married :10117	Max. :89.00	
		NA's :76	

rincome	partyid	relig
\$25000 or more:7363	Independent :4119	Protestant:10846
Not applicable:7043	Not str democrat :3690	Catholic : 5124
\$20000 - 24999:1283	Strong democrat :3490	None : 3523
\$10000 - 14999:1168	Not str republican:3032	Christian : 689
\$15000 - 19999:1048	Ind,near dem :2499	Jewish : 388
Refused : 975	Strong republican :2314	Other : 224
(Other) :2603	(Other) :2339	(Other) : 689

denom	tvhours
Not applicable :10072	Min. : 0.000
Other : 2534	1st Qu.: 1.000
No denomination : 1683	Median : 2.000
Southern baptist: 1536	Mean : 2.981
Baptist-dk which: 1457	3rd Qu.: 4.000
United methodist: 1067	Max. :24.000
(Other) : 3134	NA's :10146

However, if we summarize one column of the data set, the summary function displays many more categories.

```
summary(gss_cat$relig)
```

No answer	Don't know	Inter-nondenominational
93	15	109
Native american	Christian	Orthodox-christian
23	689	95
Moslem/islam	Other eastern	Hinduism
104	32	71
Buddhism	Other	None
147	224	3523
Jewish	Catholic	Protestant
388	5124	10846
Not applicable		
0		

Going back to the documentation for `summary()`, we can see that for a `data.frame` the default value for `maxsum` is 7, while for a factor, the default value for `maxsum` is 100. Now, this is one difference between these two summary methods, but to determine if this is the reason the outputs are different, we will need to refer to the Arguments section.

3.5 Arguments

The Arguments section lists out the name each argument to the function, the required class for that argument, and a description of what the argument does. It is probably the section I use most commonly when troubleshooting. If there are multiple functions listed in the Usage section, it is possible not every function will use every argument. You will need to refer to the Usage section to determine if a given argument applies to the function you are using.

To figure out what the `maxsum` argument actually does, we refer to the arguments section of the `summary()` help file. This tells us that `maxsum` determines the maximum number of values that should be displayed for factors. This matches the change in behavior we saw in `summary` above.

3.5.1 The ... Argument

You may see `...` as one of the arguments in the Arguments and Usage sections. This “argument” allows you to pass arguments to the function that are not explicitly listed in the function’s documentation. This is typically done for two reasons.

First, the function’s author may not want to restrict the number of arguments you can provide to the function. This is the case for functions like `sum()` and `data.frame()`. It would be very inconvenient if the R developers restricted the number of columns you could create a `data.frame` with.

```
sum(1:10, 99, 21:91, -39:45)
```

```
[1] 4385
```

```
data.frame(1:12, month.name, month.abb)
```

	X1.12	month.name	month.abb
1	1	January	Jan
2	2	February	Feb
3	3	March	Mar
4	4	April	Apr

5	5	May	May
6	6	June	Jun
7	7	July	Jul
8	8	August	Aug
9	9	September	Sep
10	10	October	Oct
11	11	November	Nov
12	12	December	Dec

This is the exception to the rule that all arguments R functions are named. If you don't want to have a set number of arguments, you can't give those arguments names.

Second, the author of a given function may want to allow other packages to expand the use cases of that function. However, the original author won't necessarily know what arguments will be useful for future R developers. So instead, they can include `...` as an argument, and then other packages can define additional arguments that will be useful to them.

The `plot` function is one example of this. If you look at the documentation for the generic `plot` function, you can see that there are only a few arguments listed, and they seemed to be geared toward X-Y plotting. This is very useful, but it is definitely not the only type of plotting you may want to do.

Many different types of objects benefit from visualization, including ones that don't neatly fit into the X-Y paradigm. It is very common that packages will define a new class of object to accomplish a particular goal and want to visualize those object in some way. Instead of creating an entirely new function to visualize their new objects, the package's creator can create a new method for `plot()`. This is useful because it standardizes the way we visualize things, and it reduces the amount of work the developer needs to do.

The `sf` package uses this functionality. The `sf` package defines the `sf` class, which stores geographical information, like the counties of North Carolina. This is the type of data we use to make maps.

```
library(sf)

#read in North Carolina data
nc = st_read(system.file("shape/nc.shp", package="sf"), quiet = TRUE)

#class of data created using sf package
class(nc)
```

```
[1] "sf"          "data.frame"
```

```
head(nc)
```

Simple feature collection with 6 features and 14 fields

Geometry type: MULTIPOLYGON

Dimension: XY

Bounding box: xmin: -81.74107 ymin: 36.07282 xmax: -75.77316 ymax: 36.58965

Geodetic CRS: NAD27

	AREA	PERIMETER	CNTY_	CNTY_ID	NAME	FIPS	FIPSNO	CRESS_ID	BIR74	SID74
1	0.114	1.442	1825	1825	Ashe	37009	37009	5	1091	1
2	0.061	1.231	1827	1827	Alleghany	37005	37005	3	487	0
3	0.143	1.630	1828	1828	Surry	37171	37171	86	3188	5
4	0.070	2.968	1831	1831	Currituck	37053	37053	27	508	1
5	0.153	2.206	1832	1832	Northampton	37131	37131	66	1421	9
6	0.097	1.670	1833	1833	Hertford	37091	37091	46	1452	7


	NWBIR74	BIR79	SID79	NWBIR79	geometry
1	10	1364	0	19	MULTIPOLYGON (((-81.47276 3...
2	10	542	3	12	MULTIPOLYGON (((-81.23989 3...
3	208	3616	6	260	MULTIPOLYGON (((-80.45634 3...
4	123	830	2	145	MULTIPOLYGON (((-76.00897 3...
5	1066	1606	3	1197	MULTIPOLYGON (((-77.21767 3...
6	954	1838	5	1237	MULTIPOLYGON (((-76.74506 3...

If we open up the R help file for plot after loading sf using `?plot`, we have the option to select documentation for the plot function in the sf package by clicking on the link labeled “plot sf object”. In the `sf::plot` help page there are many more arguments listed in the Usage and Arguments sections. In particular, we can specify the position of the key (legend), plot the data on a log scale, and add a title.

```
plot_title = 'Live Births in North Carolina from 1974-1978'

plot(nc['BIR74'], key.pos=1, logz=TRUE, main=plot_title)
```

Live Births in North Carolina from 1974–1978



3.6 Details

The this section generally contains information that is important but may not clearly fit in any of the other sections. It is often where I go looking for help when I run into a problem that I can't easily solve. Details covers the “behavior” of a function, or how a function accomplishes its task. This is especially important for functions that “behave” differently when you change the value of certain arguments. The help file for the covariance and correlation functions (`cov` and `cor`) is one example of this. The Details section describes how the functions work for each of the five options for the `use` argument. This includes telling you that one of them (“`pairwise.complete.obs`”) only works if you use `method="pearson"`. If the function in question implements a mathematical formula, the Details section may also include the formula, like in the case of `?dist`.

3.7 Value

The Value section tells you what a function will return after it completes. If you save the output of the function to a variable, this is the information that will be stored in the variable. The Value section will generally specify the class of the object as well the information the object contains. Some functions do not return a value. If this is the case, the Value section may specify that the function doesn't return anything or the Value section may be missing.

Just because something displays in your console after running a function does not mean R returned anything (see [str](#)).

3.8 Note

The Note section is a second Details section for less critical pieces of information. In the help file for the `data.frame` function, it tells us that if we need code to be compatible with versions of R earlier than R 2.4.0, the class of our `row.names` argument needs to be a character vector. While this is undoubtedly useful for a very small subset of users, it isn't relevant for most people.

3.9 Author(s)

The Authors(s) section lists the authors who wrote the function or package in question. It is more common to see for packages but there are some functions that mention specific authors as well, as is the case for `lm()`.

```
?`sf-package`
```

3.10 References

If the function implements a formula, statistical method, algorithm, procedure, or data structure developed elsewhere, the References section will include the citation for the original source. The References section for the `data.frame` function cites a textbook on the S programming language, R's precursor.

3.11 See Also

If the function whose documentation you are reviewing doesn't do exactly what you want, this is a good section to peruse. See Also contains links to the documentation files for functions that are similar to the function to the function you are investigating. This could include functions that implement a modified formula (like weighted means), or methods that are implemented for a different class like the summary method for a linear regression model (`lm`) object. It also may include methods that can be used with a given class of object, like in the case of the `lm` help file.

3.12 Examples

The Examples section is by far the most useful section for beginners. It contains example code that can be run on any computer without loading any additional packages or data. You can copy and paste it into your console and it should just work. It is a great place to start if you have absolutely no experience with a function or package, or if you have tried everything and nothing seems to work.

The Examples section can also be a good place to get code to tinker with when trying to understand how a function works. For example, if you don't understand the description of the `digits` argument in the `summary()` documentation, you can use the code in

```
#copied directly from the Examples section
summary(attenu, digits = 4)
```

event	mag	station	dist
Min. : 1.00	Min. :5.000	117 : 5	Min. : 0.50
1st Qu.: 9.00	1st Qu.:5.300	1028 : 4	1st Qu.: 11.32
Median :18.00	Median :6.100	113 : 4	Median : 23.40
Mean :14.74	Mean :6.084	112 : 3	Mean : 45.60
3rd Qu.:20.00	3rd Qu.:6.600	135 : 3	3rd Qu.: 47.55
Max. :23.00	Max. :7.700	(Other):147	Max. :370.00
		NA's : 16	

accel
Min. :0.00300
1st Qu.:0.04425
Median :0.11300
Mean :0.15422
3rd Qu.:0.21925
Max. :0.81000

We can then modify the `digits` argument to see what effect this has on the function's output.

```
summary(attenu, digits = 1)
```

event	mag	station	dist	accel
Min. : 1	Min. :5	117 : 5	Min. : 0.5	Min. :0.003
1st Qu.: 9	1st Qu.:5	1028 : 4	1st Qu.: 11.3	1st Qu.:0.044
Median :18	Median :6	113 : 4	Median : 23.4	Median :0.113
Mean :15	Mean :6	112 : 3	Mean : 45.6	Mean :0.154

```

3rd Qu.:20   3rd Qu.:7   135      :   3   3rd Qu.: 47.5   3rd Qu.:0.219
Max.      :23   Max.      :8   (Other):147   Max.      :370.0   Max.      :0.810
NA's      : 16

```

It is rare that the code in the Examples section will do exactly what you want right off the bat. That does not mean the code is not useful. The key is to use this code as a starting point, code that you know runs, and then work from there.

3.13 Further Reading

- [Getting Help with R]